

Linux问题排查

pangbin

- CPU
- memory
- I/O
- network

CPU

- %CPU
- load average

%CPU

```
top - 17:11:27 up 6:22, 3 users, load average: 0.45, 0.85, 1.07
Tasks: 119 total, 3 running, 79 sleeping, 0 stopped, 0 zombie
%Cpu0  : 98.7 us, 1.3 sy, 0.0 ni, 0.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu1  : 98.0 us, 0.7 sy, 0.0 ni, 0.3 id, 0.0 wa, 0.0 hi, 1.0 si, 0.0 st
KiB Mem : 2041132 total, 369608 free, 150688 used, 1520836 buff/cache
KiB Swap: 969960 total, 969436 free, 524 used. 1701740 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
12945	nobody	20	0	104220	11216	6412	R	98.7	0.5	0:18.20	php-fpm
12946	nobody	20	0	104220	11216	6412	R	98.3	0.5	0:18.43	php-fpm
2941	avahi	20	0	49584	5980	3208	S	1.7	0.3	11:54.00	avahi-daemon
11744	www-data	20	0	143632	7480	5616	S	0.7	0.4	0:02.82	nginx
1	root	20	0	160092	9320	6792	S	0.0	0.5	0:20.95	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kthreadd
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_wq

%CPU

- sudo **perf** record -g -p 12946
- sudo **perf** report

Samples: 9K of event 'cpu-clock', Event count (approx.): 228250000

Children	Self	Command	Shared Object	Symbol
- 99.98%	0.00%	php-fpm	[unknown]	[.] 0x2e96258d4c544155
				0x2e96258d4c544155
				__libc_start_main
- main				
- 99.62%		php_execute_script		
- zend_execute_scripts				
- 99.58%		zend_execute		
- 96.53%		execute_ex		
- 21.70%		ZEND_ASSIGN_ADD_SPEC_CV_TMPVAR_HANDLER		
- 20.30%		zend_binary_assign_op_simple_helper_SPEC_CV_TMPVAR		
8.64%		add_function		
17.38%		zif_sqrt		
8.58%		ZEND_IS_SMALLER_OR_EQUAL_SPEC_CV_CONST_HANDLER		
6.09%		__sqrt		
6.09%		ZEND_POST_INC_SPEC_CV_HANDLER		
2.44%		ZEND_FREE_SPEC_TMPVAR_HANDLER		
0.93%		ZEND_FREE_SPEC_TMPVAR_HANDLER		
0.66%		ZEND_POST_INC_SPEC_CV_HANDLER		
0.59%		ZEND_IS_SMALLER_OR_EQUAL_SPEC_CV_CONST_HANDLER		
+ 99.98%	0.00%	php-fpm	libc-2.27.so	[.] __libc_start_main
+ 99.98%	0.00%	php-fpm	php-fpm	[.] main
+ 99.62%	0.00%	php-fpm	php-fpm	[.] php_execute_script

%CPU

Overall Summary


Total Incl. Wall Time (microsec): 607 microsecs

Number of Function Calls: 1,006

[\[View Full Callgraph\]](#)

Displaying top 100 functions: Sorted by Incl. Wall Time (microsec) [[display all](#)]


Function Name	Calls	Calls%	Incl. Wall Time (microsec)	IWall%	Excl. Wall Time (microsec)	EWall%
main()	1	0.1%	607	100.0%	9	1.5%
ccccc	1	0.1%	598	98.5%	3	0.5%
bbbbbb	1	0.1%	595	98.0%	5	0.8%
aaaaaa	1	0.1%	590	97.2%	524	86.3%
sqrt	1,001	99.5%	66	10.9%	66	10.9%
xhprof_disable	1	0.1%	0	0.0%	0	0.0%



100.0%

1 call


100.0%



100.0%

1 call

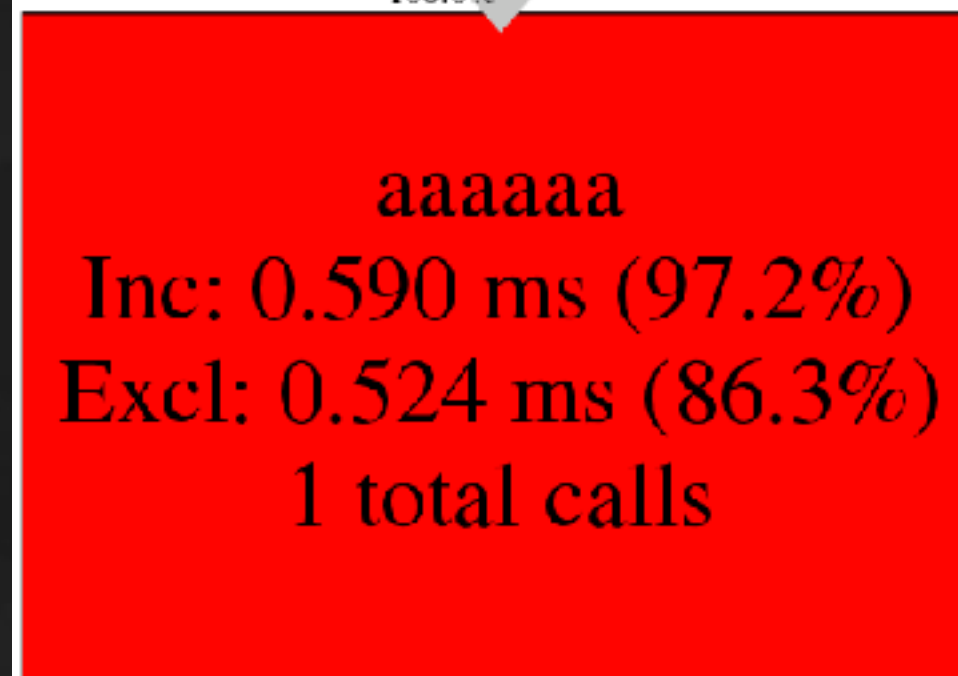
100.0%



100.0%

1 call


100.0%



100.0%

1001 calls

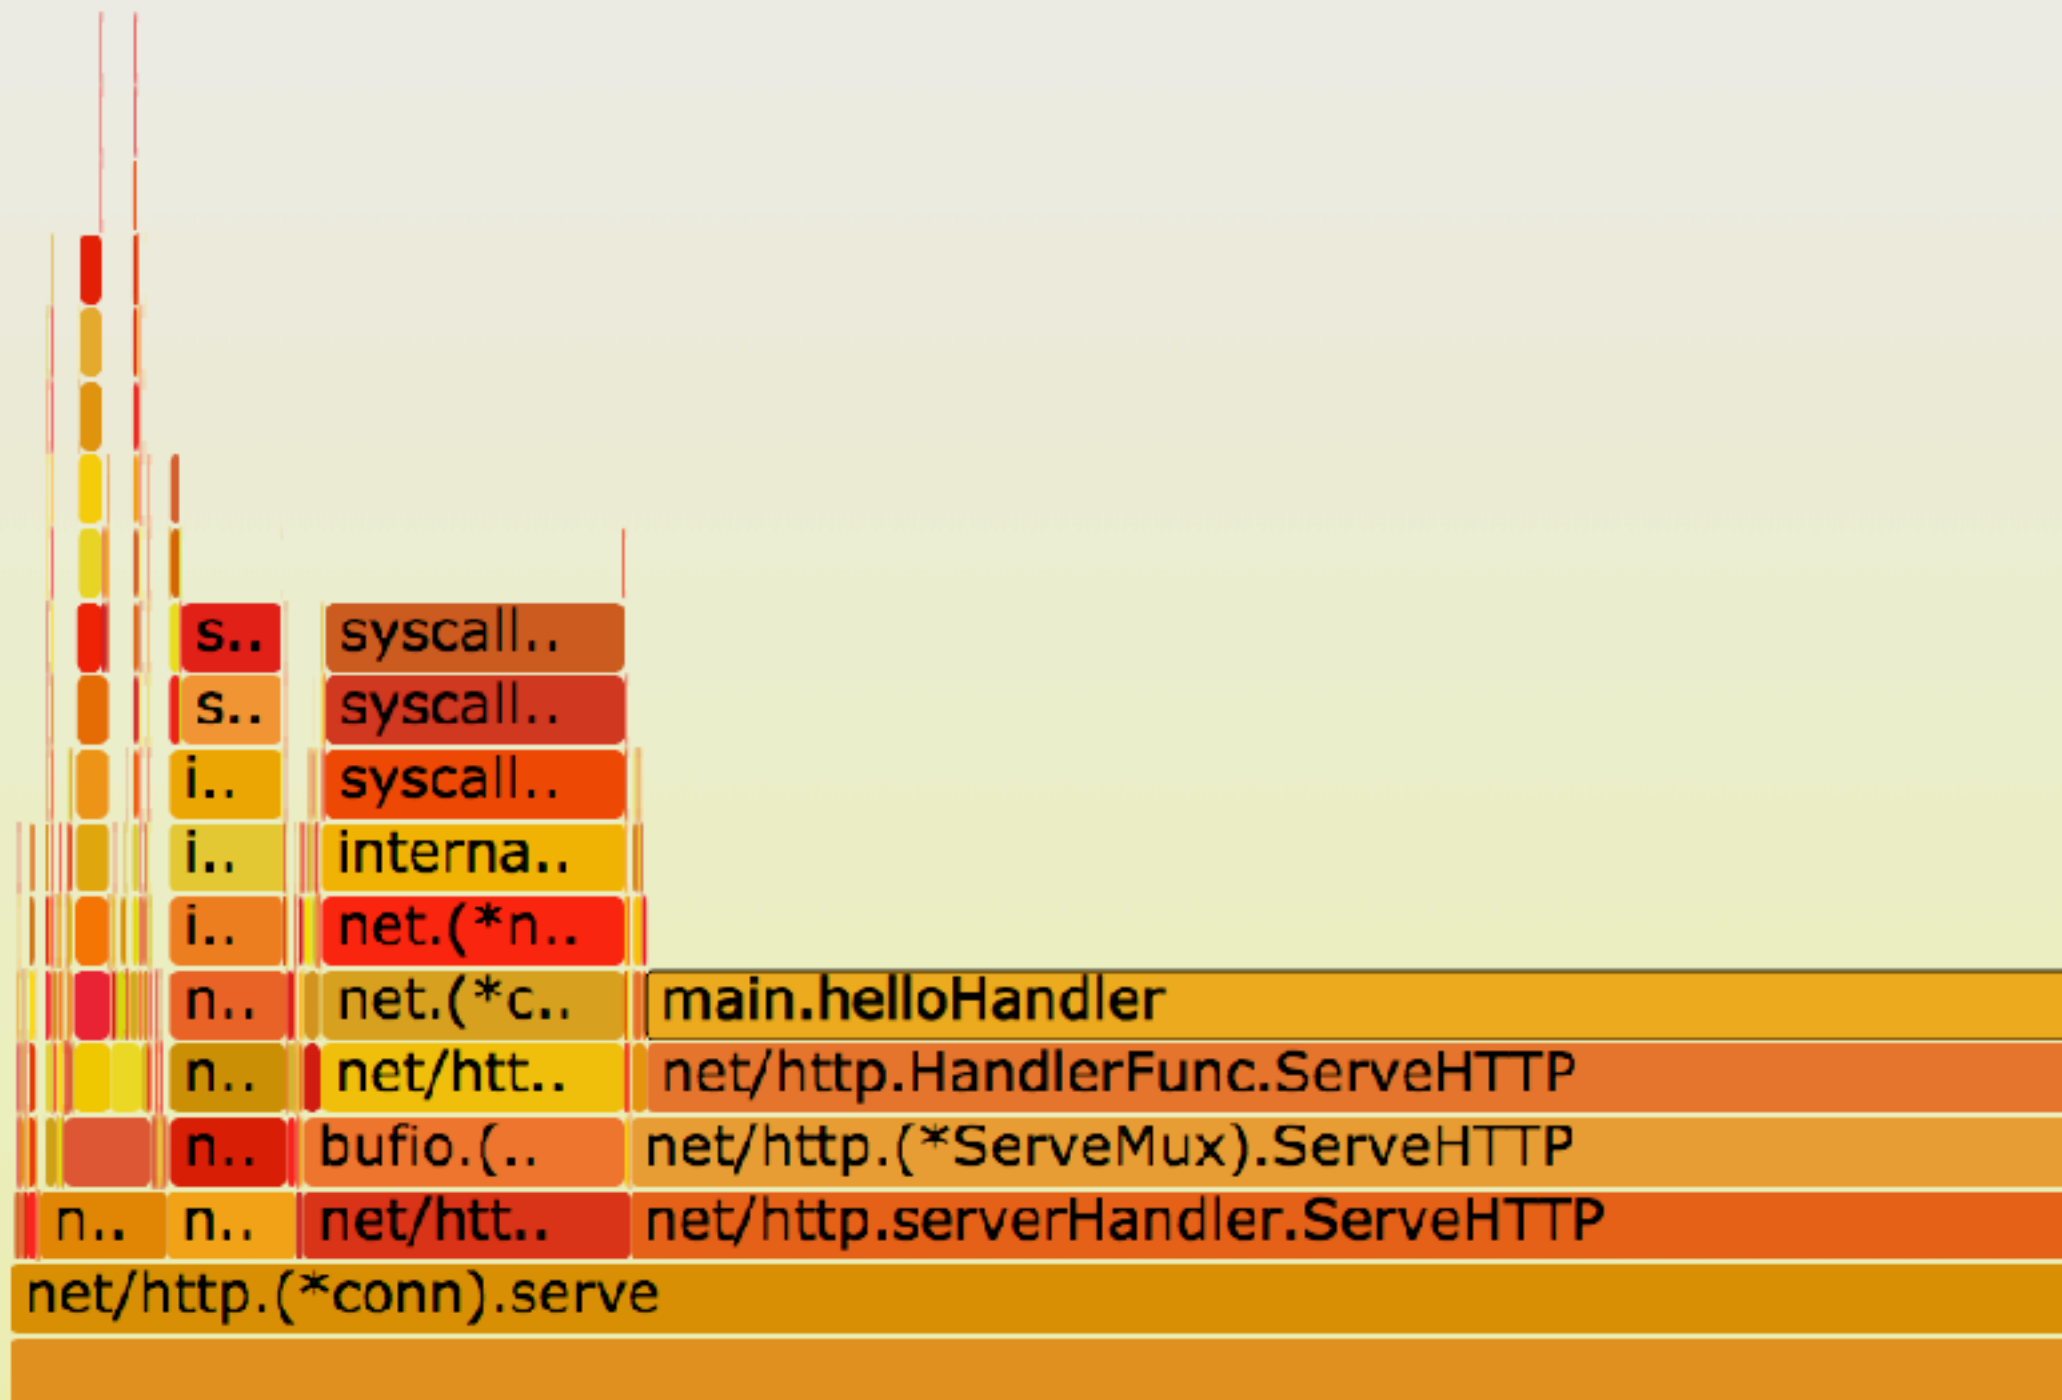
100.0%



Samples: 16K of event 'cpu-clock', Event count (approx.): 4029000000

Children	Self	Command	Shared Object	Symbol
- 95.66%	0.00%	demo	demo	[.] runtime.goexit
- runtime.goexit				
- 95.30%				net/http.(*conn).serve
- 77.97%				net/http.serverHandler.ServeHTTP
- net/http.(*ServeMux).ServeHTTP				
- 77.77%				net/http.HandlerFunc.ServeHTTP
- 77.75%				main.helloHandler
- 13.39%				swaps_restore_regs_and_return_to_usermode
prepare_exit_to_usermode				
- exit_to_usermode_loop				
- 2.23%				schedule
__schedule				
finish_task_switch				
+ 0.66%				ret_from_intr
+ 0.65%				fmt.Sprintf
+ 7.14%				net/http.(*conn).serve.func1
+ 6.29%				net/http.(*response).finishRequest
+ 2.67%				net/http.(*conn).readRequest
+ 0.73%				net.(*TCPAddr).String
+ 95.30%	0.04%	demo	demo	[.] net/http.(*conn).serve
+ 77.97%	0.01%	demo	demo	[.] net/http.serverHandler.ServeHTTP
+ 77.97%	0.02%	demo	demo	[.] net/http.(*ServeMux).ServeHTTP
+ 77.77%	0.00%	demo	demo	[.] net/http.HandlerFunc.ServeHTTP
77.75%	62.65%	!	!	[.] main.helloHandler

FlameGraph



Function: `main.helloHandler` (4,860 samples, 85.01%)

Compiler optimization

- PHP: Makefile
CFLAGS_CLEAN = -I/usr/include -g -O0 -
fvisibility=hidden -DZEND_SIGNALS \$(PROF_FLAGS)
- Go: go build -gcflags '-N -l' demo.go

load average

- define: System load averages is the average number of processes that are either in a runnable or uninterruptable state
- high?: load average: 4.90, 1.66, 0.59
- high?: how to do

memory

- memory leak
- large memory usage

memory leak

- memleak
- valgrind
- pprof
- xhprof
-


memory leak

```
1 <?php
2
3 class Foo {
4     public $var = '3.14159265359';
5 }
6
7 $baseMemory = memory_get_usage();
8
9 for ($i = 0; $i <= 100000; $i++) {
10     $a = new Foo;
11     $a->self = $a;
12     if ($i % 500 === 0) {
13         echo sprintf('%8d: ', $i), memory_get_usage() - $baseMemory, "\n";
14     }
15 }
16
17
18
19 █
```

~

~

memory leak



large memory usage

```
top - 17:44:22 up 7:03, 5 users, load average: 0.09, 0.04, 0.01
Tasks: 129 total, 1 running, 93 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.3 us, 0.2 sy, 0.0 ni, 99.3 id, 0.0 wa, 0.0 hi, 0.2 si, 0.0 st
KiB Mem : 2041132 total, 297968 free, 582152 used, 1161012 buff/cache
KiB Swap: 969960 total, 969960 free, 0 used. 1268772 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
4753	pangbin	20	0	529584	451468	10516	S	1.7	22.1	0:08.36	php
590	root	20	0	777280	22440	12256	S	0.0	1.1	0:02.15	snaped
220	root	19	-1	103072	21696	20892	S	0.0	1.1	0:00.33	systemd-journal
589	root	20	0	177624	17648	9756	S	0.0	0.9	0:00.19	networkd-dispat
721	root	20	0	297088	16940	13796	S	0.0	0.8	0:01.35	php-fpm7.2
575	root	20	0	565620	16884	14092	S	0.0	0.8	0:00.79	NetworkManager
1466	root	20	0	373672	13996	12300	S	0.0	0.7	0:00.15	packagekitd
925	whoopsie	20	0	388416	12784	11160	S	0.0	0.6	0:00.03	whoopsie
2015	root	20	0	303656	11008	9588	S	0.0	0.5	0:00.00	cups-browsed
756	www-data	20	0	299728	10872	7360	S	0.0	0.5	2:55.45	php-fpm7.2
755	www-data	20	0	299728	10624	7116	S	0.0	0.5	2:56.09	php-fpm7.2
1	root	20	0	225444	9232	6784	S	0.0	0.5	0:02.52	systemd
582	root	20	0	427260	8948	7588	S	0.0	0.4	0:00.06	ModemManager
1350	pangbin	20	0	76724	7952	6772	S	0.0	0.4	0:00.03	systemd

php-memory-profiler

```
<?php
```

```
memprof_enable();
```

```
//your php code
```

```
//maybe use max_execution_time and
```

```
//register_shutdown_function
```

```
var_dump(memprof_dump_array());
```

```

{
  "memory_size": 36456,
  "blocks_count": 302,
  "memory_size_inclusive": 553405,
  "blocks_count_inclusive": 606,
  "calls": 1,
  "called_functions": {
    "unknown": {
      "memory_size": 0,
      "blocks_count": 0,
      "memory_size_inclusive": 0,
      "blocks_count_inclusive": 0,
      "calls": 200,
      "called_functions": []
    },
    "UserCenter_User::GetUserInfo": {
      "memory_size": 0,
      "blocks_count": 0,
      "memory_size_inclusive": 0,
      "blocks_count_inclusive": 0,
      "calls": 200,
      "called_functions": {
        "str_repeat": {
          "memory_size": 0,
          "blocks_count": 0,
          "memory_size_inclusive": 0,
          "blocks_count_inclusive": 0,
          "calls": 20000,
          "called_functions": []
        }
      }
    }
  },
  "str_repeat": {
    "memory_size": 516600,
    "blocks_count": 300,
    "memory_size_inclusive": 516600,
    "blocks_count_inclusive": 300,
    "calls": 300,
    "called_functions": []
  },
  "memprof_dump_array": {"memory_size": 349...}
}

```

- memory_size
- blocks_count
- memory_size_inclusive
- blocks_count_inclusive
- calls
- called_functions

I/O

```
top - 20:42:36 up 32 min,  2 users,  load average: 1.48, 1.63, 1.35
Tasks: 143 total,   2 running, 106 sleeping,   0 stopped,   0 zombie
%Cpu(s):  1.6 us,  3.8 sy,   0.0 ni, 56.2 id, 35.4 wa,   0.0 hi,  2.9 si,   0.0 st
KiB Mem : 2041132 total, 1045988 free,  320396 used,  674748 buff/cache
KiB Swap:  969960 total,  969960 free,    0 used. 1568636 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1077	redis	20	0	140472	8176	2768	D	10.3	0.4	1:37.08	redis-server
1666	root	20	0	61388	14756	7216	S	5.0	0.7	0:05.26	iotop
184	root	20	0	0	0	0	R	3.3	0.0	0:34.38	jbd2/sda1-8
164	root	0	-20	0	0	0	I	2.0	0.0	0:25.02	kworker/0:1H
1	root	20	0	160004	9392	6796	S	0.0	0.5	0:01.89	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_wq
7	root	20	0	0	0	0	S	0.0	0.0	0:01.54	ksoftirqd/0
8	root	20	0	0	0	0	I	0.0	0.0	0:01.31	rcu_sched
9	root	20	0	0	0	0	I	0.0	0.0	0:00.00	rcu_bh
10	root	rt	0	0	0	0	S	0.0	0.0	0:00.00	migration/0

I/O

Total DISK READ :		0.00 B/s	Total DISK WRITE :		1007.91 K/s		
Actual DISK READ:		0.00 B/s	Actual DISK WRITE:		3.95 M/s		
TID	PRIO	USER	DISK READ	DISK WRITE	SWAPIN	IO>	COMMAND
184	be/3	root	0.00 B/s	0.00 B/s	0.00 %	76.83 %	[jbd2/sda1-8]
1077	be/4	redis	0.00 B/s	1007.91 K/s	0.00 %	3.75 %	redis-server *:6379
1164	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.02 %	[kworker/u4:0]
1	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	init splash
2	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kthreadd]
4	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kworker/0:0H]
6	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[mm_percpu_wq]
7	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[ksoftirqd/0]
8	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[rcu_sched]
9	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[rcu_bh]
10	rt/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[migration/0]
11	rt/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[watchdog/0]
12	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[cpuhp/0]
13	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[cpuhp/1]
14	rt/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[watchdog/1]
15	rt/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[migration/1]
16	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[ksoftirqd/1]
18	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kworker/1:0H]
19	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kdevtmpfs]
20	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[netns]
21	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[rcu_tasks_kthre]
22	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kauditd]
24	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[khungtaskd]
25	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[oom_reaper]

I/O

```
pangbin@pangbin-VirtualBox:~$ sudo lsof -p 1077
```

COMMAND	PID	USER	FD	TYPE	DEVICE	SIZE/OFF	NODE	NAME
redis-ser	1077	redis	cwd	DIR	8,1	4096	942111	/var/lib/redis
redis-ser	1077	redis	rtd	DIR	8,1	4096	2	/
redis-ser	1077	redis	txt	REG	8,1	1085080	399689	/usr/bin/redis-check-rdb
redis-ser	1077	redis	mem	REG	8,1	10281936	400281	/usr/lib/locale/locale-archive
redis-ser	1077	redis	mem	REG	8,1	2030544	137020	/lib/x86_64-linux-gnu/libc-2.27.so
redis-ser	1077	redis	mem	REG	8,1	144976	137153	/lib/x86_64-linux-gnu/libpthread-2.27.so
redis-ser	1077	redis	mem	REG	8,1	215656	399685	/usr/lib/x86_64-linux-gnu/libjemalloc.so.1
redis-ser	1077	redis	mem	REG	8,1	14560	137043	/lib/x86_64-linux-gnu/libdl-2.27.so
redis-ser	1077	redis	mem	REG	8,1	1700792	137083	/lib/x86_64-linux-gnu/libm-2.27.so
redis-ser	1077	redis	mem	REG	8,1	170960	136992	/lib/x86_64-linux-gnu/ld-2.27.so
redis-ser	1077	redis	0u	CHR	1,3	0t0	9	/dev/null
redis-ser	1077	redis	1u	CHR	1,3	0t0	9	/dev/null
redis-ser	1077	redis	2u	CHR	1,3	0t0	9	/dev/null
redis-ser	1077	redis	3r	FIFO	0,12	0t0	23258	pipe
redis-ser	1077	redis	4w	FIFO	0,12	0t0	23258	pipe
redis-ser	1077	redis	5u	a_inode	0,13	0	9590	[eventpoll]
redis-ser	1077	redis	6u	IPv6	24487	0t0	TCP	*:6379 (LISTEN)
redis-ser	1077	redis	7u	IPv4	24488	0t0	TCP	*:6379 (LISTEN)
redis-ser	1077	redis	8w	REG	8,1	45979035	942147	/var/lib/redis/appendonly.aof
redis-ser	1077	redis	9u	IPv4	27546	0t0	TCP	pangbin-VirtualBox:6379->10.106.10.202:62345 (ESTABLISHED)

```
[pid 1077] 20:50:43.466249 epoll_wait(5, [{EPOLLIN, {u32=9, u64=9}}], 10128, 39) = 1 <0.000596>
[pid 1077] 20:50:43.467412 read(9, "*3\r\n$3\r\nSET\r\n$16\r\nkey:0000153724"... , 16384) = 45 <0.000869>
[pid 1077] 20:50:43.468761 read(3, 0x7fff15267857, 1) = -1 EAGAIN (Resource temporarily unavailable) <0.000404>
[pid 1077] 20:50:43.469858 write(8, "*3\r\n$3\r\nSET\r\n$16\r\nkey:0000153724"... , 45) = 45 <0.000481>
[pid 1077] 20:50:43.470813 fdatasync(8) = 0 <0.004323>
[pid 1077] 20:50:43.475391 write(9, "+OK\r\n", 5) = 5 <0.000953>
[pid 1077] 20:50:43.476894 epoll_wait(5, [{EPOLLIN, {u32=9, u64=9}}], 10128, 28) = 1 <0.000407>
```

network

- domain
- packet loss
- network traffic


```
pangbin@pangbin-VirtualBox:~$ traceroute example.com
traceroute to example.com (93.184.216.34), 30 hops max, 60 byte packets
 1  _gateway (192.168.119.254)  8.721 ms  8.672 ms  8.942 ms
 2  10.0.246.1 (10.0.246.1)  8.779 ms  9.311 ms  8.515 ms
 3  CVRouter.lan (10.106.144.1)  8.629 ms  8.609 ms  8.573 ms
 4  103.206.191.1 (103.206.191.1)  8.547 ms  8.524 ms  8.501 ms
 5  100.101.71.1 (100.101.71.1)  13.990 ms  *  *
 6  100.101.66.1 (100.101.66.1)  46.994 ms  39.488 ms  40.064 ms
 7  100.101.66.22 (100.101.66.22)  78.886 ms  80.786 ms  79.577 ms
 8  10.16.13.141 (10.16.13.141)  108.172 ms  75.100 ms  75.083 ms
 9  10.16.208.54 (10.16.208.54)  64.688 ms  64.665 ms  45.186 ms
10  ce-0-12-0-3.r03.tkokhk01.hk.bb.gin.ntt.net (203.131.240.241)  47.241 ms  47.233 ms
11  ae-11.r24.tkokhk01.hk.bb.gin.ntt.net (129.250.6.99)  48.434 ms  48.265 ms  49.107 ms
12  ae-12.r30.tokyjp05.jp.bb.gin.ntt.net (129.250.2.50)  89.742 ms  96.745 ms  97.445 ms
13  ae-4.r23.snjsca04.us.bb.gin.ntt.net (129.250.5.78)  200.103 ms  204.808 ms  204.984 ms
14  ae-45.r01.snjsca04.us.bb.gin.ntt.net (129.250.3.175)  214.820 ms  198.385 ms  204.760 ms
15  ae-0.edgecast-networks.snjsca04.us.bb.gin.ntt.net (129.250.193.134)  204.243 ms  199.959 ms
    93.250)  204.524 ms
16  152.195.84.133 (152.195.84.133)  200.338 ms  152.195.85.133 (152.195.85.133)  199.959 ms
17  93.184.216.34 (93.184.216.34)  201.285 ms  204.090 ms  204.944 ms
18  93.184.216.34 (93.184.216.34)  204.604 ms  200.445 ms  201.615 ms
    100.101.71.1 (100.101.71.1)  13.990 ms  *  *
```


domain

- nslookup
- dig
- ping

packet loss

```
pangbin@pangbin-VirtualBox:~$ sar -n EDEV 1 5
```

```
Linux 4.15.0-20-generic (pangbin-VirtualBox)
```

2019年06月10日

_x86_64_

(2 CPU)

[illegible]

packet

```
pangbin@pangbin-VirtualBox:~$ sar -n DEV 1 5
```

```
Linux 4.15.0-20-generic (pangbin-VirtualBox)
```

2019年06月10日

_x86_64_

(2 CPU)

[illegible]

socket

Tcp:

- 6 active connection openings
- 6 passive connection openings
- 0 failed connection attempts
- 4 connection resets received
- 2 connections established
- 537926 segments received
- 434182 segments sent out
- 0 segments retransmitted
- 0 bad segments received
- 4 resets sent

Udp:

- 262 packets received
- 3 packets to unknown port received
- 0 packet receive errors
- 224 packets sent
- 0 receive buffer errors
- 0 send buffer errors
- IgnoredMulti: 81

UdpLite:

TcpExt:

- 4 delayed acks sent
- Quick ack mode was activated 9 times
- 430 packet headers predicted

network traffic

No.	Time	Source	Destination	Protocol	Length	Info
477	3.631282	10.106.10.202	10.106.10.204	TCP	78	56950 → 80 [SYN] Seq=0 Win=655
478	3.631547	10.106.10.204	10.106.10.202	TCP	74	80 → 56950 [SYN, ACK] Seq=0 Ac
479	3.631584	10.106.10.202	10.106.10.204	TCP	66	56950 → 80 [ACK] Seq=1 Ack=1 W
480	3.631684	10.106.10.202	10.106.10.204	HTTP	143	GET / HTTP/1.1
481	3.631885	10.106.10.204	10.106.10.202	TCP	66	80 → 56950 [ACK] Seq=1 Ack=78
482	3.632112	10.106.10.204	10.106.10.202	HTTP	925	HTTP/1.1 200 OK (text/html)
483	3.632136	10.106.10.202	10.106.10.204	TCP	66	56950 → 80 [ACK] Seq=78 Ack=86
484	3.632292	10.106.10.202	10.106.10.204	TCP	66	56950 → 80 [FIN, ACK] Seq=78 A
485	3.632498	10.106.10.204	10.106.10.202	TCP	66	80 → 56950 [FIN, ACK] Seq=860
486	3.632547	10.106.10.202	10.106.10.204	TCP	66	56950 → 80 [ACK] Seq=79 Ack=86

- ▶ Frame 480: 143 bytes on wire (1144 bits), 143 bytes captured (1144 bits) on interface 0
- ▶ Ethernet II, Src: Apple_ba:35:30 (3c:15:c2:ba:35:30), Dst: PcsCompu_60:82:e0 (08:00:27:60:82:e0)
- ▶ Internet Protocol Version 4, Src: 10.106.10.202, Dst: 10.106.10.204
- ▶ Transmission Control Protocol, Src Port: 56950, Dst Port: 80, Seq: 1, Ack: 1, Len: 77
- ▼ Hypertext Transfer Protocol
 - ▶ GET / HTTP/1.1\r\n
 - Host: 10.106.10.204\r\n
 - User-Agent: curl/7.63.0\r\n
 - Accept: */*\r\n
 - \r\n
 - [Full request URI: http://10.106.10.204/]
 - [HTTP request 1/1]
 - [Response in frame: 482]



thanks

- <https://www.ostechnix.com/the-top-command-tutorial-with-examples-for-beginners/>
- <https://github.com/BitOne/php-meminfo>
- <https://medium.com/@malith.jayasinghe/network-monitoring-using-sar-37bab6ce9f68>